Amazon Product Catalogue

Pricing, Visibility & Review Drivers

Data Visualization & Analytics

Capstone 2 — Final Project Report

Sector: E-commerce / Retail

Institute: Newton School of Technology

Team Members:
Jagadish Ishwar Patil
Manthan Subhash Ziman
Uzair Ahmed Shah
Raman Pandey
Aashish Jha

2. Executive Summary

Problem. Amazon sellers and category managers face a noisy, crowded catalogue: 42K+ SKUs visible across browse pages, but only a handful achieve best-seller status or sustained review volume. Without a quantitative read on what drives visibility and trust signals, pricing and promo decisions are guesswork.

Approach. Starting from a raw 42,675-row Amazon scrape (uncleaned variant), we built a Python ETL pipeline that parsed currency, ratings, review-shorthand, and badge text into clean numeric/boolean columns; deduplicated by title (Amazon's search returns the same SKU on multiple pages — 86% duplicates removed); engineered features for selling_price, discount_pct, brand (heuristic), price_tier, rating_bucket, and log-transforms of skewed counts. The cleaned **5,971-SKU × 21-column** dataset (zero NaN values across all analytical columns) feeds a Tableau Public dashboard and a statistical analysis covering correlation, one-way ANOVA, Welch's t-test, OLS regression, and KMeans segmentation.

Key Insights.

1. **Best-sellers are cheap, not premium** — median price for best-seller-badged SKUs is **$26 vs $75** for the rest (n=160 vs 5,811), with **11× higher median review count** (10,498 vs 949).
2. **Discount depth has no relationship with rating** (ANOVA F=0.24, p=0.92) — discounting does not earn or destroy customer satisfaction in this catalogue.
3. **Sponsored slots underperform on ratings** (4.38 vs 4.46, Welch t=-3.75, p<0.001) and have 7× lower median review volume — sponsored placement is a discoverability tool for newer / less-mature SKUs, not a quality signal.
4. **Price is negatively correlated with engagement** — log(reviews) ↔ selling_price r=-0.36, log(bought_last_month) ↔ selling_price r=-0.43. Cheaper SKUs accumulate engagement faster.

Key Recommendations.

1. **For sellers seeking volume, target the Budget/Mid tier (≤$50) where 80% of SKUs sit at Excellent ratings.** Premium-tier SKUs (>$150) drop to 43% Excellent and earn 4× fewer reviews.
2. **Stop using discount depth as the primary lever for rating improvement** — it has no measurable effect. Reinvest that margin in listing quality and keyword optimisation instead.
3. **Treat sponsored ad spend as a launch tool, not a flagship channel** — sponsored SKUs both rate and engage worse than organic peers; use ads to seed early reviews on new SKUs, then let organic ranking compound.

3. Sector & Business Context

Sector overview. E-commerce on Amazon is a high-volume, low-information environment for sellers: ~600M products globally, with ~2M active third-party sellers competing on a small set of decision-relevant signals (price, rating, review count, badges, ad slot). India's Amazon ecosystem alone generated $32B+ GMV in 2024 with 14M+ sellers globally on the platform.

Decision-maker. This project serves a **third-party Amazon seller / category manager** who must allocate finite resources across:

- new-SKU launches (which categories, which price points?),
- promotional spend (ads vs coupons vs deeper discounts),
- portfolio pruning (which slow movers should be retired?).

Why this problem. Sellers default to gut-feel rules — "discount more to get reviews", "sponsored ads always pay off", "premium = quality halo". This study tests those rules of thumb against a real cross-sectional snapshot of the catalogue.

Business value. A 1% improvement in best-seller-rate within an active 1,000-SKU portfolio is worth low-six-figure annual revenue at typical Amazon AOVs. A correct redirection of ad spend (currently typically 8–12% of revenue for many sellers) toward listing quality compounds over the SKU lifecycle.

4. Problem Statement & Objectives

Formal problem definition. Given a snapshot of 42K+ Amazon listings, identify the pricing, discount, and badge characteristics that distinguish high-engagement SKUs (review volume, best-seller status, recent purchase velocity), and translate those characteristics into actionable portfolio decisions.

Project scope (in).

- Cross-sectional analysis of a single ~2025-08 snapshot.
- Catalogue-wide patterns (no individual-SKU recommendations).
- USD-denominated price analysis.
- Heuristic brand attribution (first token of title).

Out of scope.

- Time-series / forecasting (single snapshot). Category-
- level analysis (no category column in source). Causal
- inference (this is observational data).
- Inventory / margin / fulfilment cost data (not in source).

Success criteria.

- Reproducible Python ETL pipeline producing ≥5,000 clean rows × 8 cols (■ 5,971 × 21, zero NaN).
- Statistical tests with reported test statistics, p-values, and effect direction.
- Tableau Public dashboard with interactive filters answering ≥3 stakeholder questions.
- 8–12 written, decision-language insights and 3–5 prioritised recommendations.

5. Data Description

Source. Kaggle — Amazon Products Sales Dataset (42K Items, 2025) by ikramshah512. URL: https://www.kaggle.com/datasets/ikramshah512/amazon-products-sales-dataset-42k-items-2025 File used: `amazon_products_sales_data_uncleaned.csv (~36.7 MB)`.

Structure. 42,675 rows × 16 columns. Single snapshot collected ~2025-08-21. Currency = USD throughout (`listed_price` uses `$`).

Key raw columns.

Column	Type	Cleaning needed
`title`	free text	None — used as primary key for dedupe; brand parsed from first token
`rating`	string `"4.6 out of 5 stars"`	Regex extract → float, clip 0–5
`number_of_reviews`	string `"1,234"` / `"6K+"`	Strip commas, expand K/M shorthand
`bought_in_last_month`	string `"300+ bought in past month"`	Regex extract, expand shorthand
`current/discounted_price`	numeric-string	Float coerce; 27.5% missing
`listed_price`	string `"$159.00"`	Strip $ and comma → float
`is_best_seller`	string badge text	Boolean: contains "best" or "amazon's choice"
`is_sponsored`	string	Boolean: contains "sponsored"
`is_couponed`	string	Extract % if present; else 0 / NaN for $-coupons
`buy_box_availability`	string	Boolean: "Add to cart" / "Buy now"
`sustainability_badges`	string	Boolean: any non-null value (92% null)
`image_url`, `product_url`, `delivery_details`, `price_on_variant`	strings	Dropped — non-analytical or noise
`collected_at`	string	Parsed to datetime

Compliance with Capstone 2 input-dataset rules. The Capstone spec requires the source dataset (what we drop into data/raw/) to be raw or minimally processed, with row-level records, real missing values, and no pre-built dashboard or feature-engineered version. The chosen file satisfies every clause:

Spec rule	Evidence in `amazon_products_sales_data_uncleaned.csv`
Raw / minimally processed	File is named `_uncleaned`. Prices are stored as `"$159.00"`, ratings as `"4.6 out of 5 stars"`, review counts as `"6K+"`, coupons as `"Save 15% with coupon"`. None of it is analysis-ready.
Row-level records available	42,675 rows, one row per SKU listing
Missing values & inconsistencies present	`current_discounted_price 27.5% NaN`, `listed_price 54% NaN`, `sustainability_badges 92%` NaN, buy_box_availability 34% NaN; 24,912 duplicate listings across paginated search results
No pre-built dashboard or feature-engineered version	`No brand`, `no price_tier`, `no discount_pct`, `no` log-transforms — every engineered column in our final dataset is produced by our pipeline
Volume — ≥ 5,000 rows, ≥ 8 analytical columns	42,675 raw rows, 16 raw columns; final 5,971 × 21

The cleaning, deduping, imputation, and feature engineering described in §6 are exactly what the Capstone requires us to demonstrate against this raw input — they are not violations of the input rule.

Limitations & known biases.

- **No category taxonomy in source** analysis is product-level only; brand is a heuristic from the first title token.
- **Single snapshot** no time series, so trend analysis is not possible.
- **Search-result bias** the scrape captures pages a logged-out user sees, skewed toward popular / sponsored placements.
- **Title-based dedupe is aggressive** Amazon search repeats listings across pagination; we kept first occurrence per title (~86% dropped).

6. Data Cleaning & ETL Pipeline

All cleaning is implemented in `scripts/etl_pipeline.py` and mirrored cell-by-cell in `notebooks/02_cleaning.ipynb` (Capstone requirement: Python + Jupyter).

Pipeline stages

Step	Action	Result
1	Load raw CSV	42,675 × 16
2	Standardise column names → `snake_case` Parse string-encoded numerics (rating, review counts, currency, coupon %)	rename `current/discounted_price` → `current_discounted_price`
3	Cap rating to [0, 5]; coupon % to [0, 100]	dtype changes (see below)
4	Convert badge fields → booleans (is_best_seller,	physical-bound enforcement
5	`is_sponsored, has_buy_box, has_coupon, has_sustainability_badge)` Drop rows with empty title or no price; dedupe by title	analysis-friendly types
6	Engineer	42,675 → 6,014 rows (-85.9%)
7	`selling_price = current_discounted_price.fillna(listed_price)`	unified price field
8	`Engineer discount_pct, is_discounted, brand, price_tier, rating_bucket, log_reviews, log_bought_last_month`	21 final analytical columns
9	Impute remaining NaN per the policy below; drop SKUs with no rating (43 rows)	6,014 → **5,971** rows
10	Drop noise columns (URLs, image, raw badge strings)	clean output
11	`Write data/processed/amazon_products_clean.csv`	**zero NaN** in every analytical column

Missing-value handling

The pipeline guarantees a fully-imputed final dataset. Every column written to amazon_products_clean.csv has zero NaN values; this is enforced by a sanity check at the end of engineer_features() that logs a WARNING if any NaN remains.

	Raw % missing	Treatment
`sustainability_badges`	92.0%	Replaced with boolean `has_sustainability_badge; raw` column dropped
`buy_box_availability`	34.3%	Replaced with boolean has_buy_box; raw column dropped
`listed_price`	54.0%	`Filled with selling_price` (no-discount assumption) **after** discount_pct is computed
`current_discounted_price`	27.5%	Filled with selling_price after discount is computed
`bought_in_last_month`	≈5%	Filled with 0 (absence ≡ no recent purchase activity)
`number_of_reviews`	<1%	Filled with 0 (absence ≡ no review data)
`coupon_pct`	small	Filled with 0 (NaN means a non-percentage $-coupon, presence is preserved by has_coupon)
`rating`	0.7%	**Dropped** (43 rows). Median-imputing 4.5 would distort downstream tests

Outlier handling

- **Price** is heavy-tailed (max $4,699). We **flag** but do not drop (they are real luxury items: cameras, laptops). Log-transform used for any modelling.
- **Reviews** is heavy-tailed (max ≈ 866K). Log-transform used everywhere.

Feature engineering decisions

Eight new analytical columns are derived from the raw schema:

New column	Source	Logic
`selling_price`	`current_discounted_price, listed_price`	discounted price if present, else listed
`discount_pct`	`listed_price, current_discounted_price`	(listed − discounted) / listed × 100
`is_discounted`	`discount_pct`	True if discount_pct > 0
`coupon_pct`	`is_couponed (text)`	regex extract of %; non-percent coupons → 0
`has_coupon`	`coupon_pct`	True if coupon_pct > 0
`brand`	`title`	first cleaned token (heuristic)
`price_tier`	`selling_price`	Budget ≤$15 / Mid ≤$50 / Premium ≤$150 / Luxury >$150
`rating_bucket`	`rating`	Poor ≤2.9 / Average ≤3.9 / Good ≤4.4 / Excellent >4.4
`log_reviews, log_bought_last_month`	counts	log1p to tame heavy tails

- brand **is a heuristic** — first cleaned token of title. Acceptable for ranking the top 10 brands (HP, Amazon, Samsung, Canon, Epson, Ugreen, Asus, Lenovo, Brother, Garmin all appear correctly), but brittle for tail brands (e.g. "Complete" from "Complete Protect" warranty plans).

- **Price tiers in USD**: Budget ≤$15 / Mid ≤$50 / Premium ≤$150 / Luxury >$150. Tuned to roughly equal-size buckets given the catalogue distribution.

Assumptions made during cleaning

1. A row with no title is unanalysable — drop.
2. A row with no price (both current_discounted_price AND listed_price NaN) is unanalysable — drop.
3. Two rows with identical titles are duplicate listings of the same SKU on different result pages — keep the first occurrence.
4. NaN rating cannot be safely median-imputed without distorting the ANOVA / t-test / regression that depend on it — drop those 43 rows.
5. NaN engagement counts (number_of_reviews, bought_in_last_month) mean "no recent activity / no review" — fill with 0.
6. NaN coupon% means a $-coupon (not %); fill with 0 but preserve presence via has_coupon.
7. Heavy-tailed price/review distributions are log-transformed for any modelling; raw values retained for descriptive stats.

Reference. All cleaning logic is implemented in scripts/etl_pipeline.py and mirrored cell-by-cell in notebooks/02_cleaning.ipynb (Capstone requirement). The pipeline is reproducible — running python scripts/etl_pipeline.py from project root regenerates `data/processed/amazon_products_clean.csv deterministically.`

7. KPI & Metric Framework

KPI	Value	Formula	Why it matters	
Catalogue Size	**5,971**	`count(rows)`	Inventory scope after dedupe + rating imputation	
Avg Selling Price	**$180.81**	`mean(selling_price)`	Pricing positioning	
Median Selling Price	**$71.95**	`median(selling_price)`	Robust price level (tail-resistant)	
Weighted Avg Rating	**4.578**	`Σ(rating × reviews)/ Σreviews`	Real customer voice (weights stronger SKUs)	
Avg Rating (unweighted)	4.453	`mean(rating)`	Catalogue-wide quality baseline	
Total Reviews	**37.9 M**	`sum(number_of_reviews)`	Aggregate trust signal	
% SKUs Discounted	**46.12%**	`mean(is_discounted)`		
Avg Discount when on sale	**21.1%**	`mean(discount_pct \`	Promo aggressiveness	Discount depth
% Sponsored	**4.69%**	`mean(is_sponsored)`	discount > 0)`	
% Best Seller / Amazon's Choice	**2.68%**	`mean(is_best_seller)`	Ad-spend share Catalogue strength	
% With Coupon	**1.59%**	`mean(has_coupon)`	Coupon usage	

KPI	Value	Formula	Why it matters	
% Buy Box available	**91.41%**	`mean(has_buy_box)`	Listing health	
% Sustainability badge	**8.94%**	`mean(has_sustainability_badge)`	Sustainability participation	
Unique brands (heuristic)	570	`nunique(brand)`	Catalogue diversity	

The KPI table is computed in `notebooks/05_final_load_prep.ipynb` and saved to `data/processed/amazon_kpis.csv for the Tableau dashboard.`

Mapping KPIs to project objectives

Every KPI ties to one of the four objectives stated in §4:

Project objective	KPIs that measure it
O1 — Quantify pricing positioning across the catalogue	Avg / Median Selling Price, % SKUs Discounted, Avg Discount when on sale, price-tier shares
O2 — Surface visibility & trust signals (badges, ad slot, buy-box)	% Best Seller, % Sponsored, % Buy Box, % With Coupon, % Sustainability badge
O3 — Establish a customer-quality baseline	Weighted Avg Rating, Avg Rating, Total Reviews, rating-bucket shares
O4 — Identify portfolio-level structure (brand mix, segment density)	Catalogue Size, Unique brands, top-10 brand share, KMeans cluster sizes (§9.5)

Recommendations in §12 are then framed as moves on these KPIs: e.g. R1 (discount cap) targets Avg Discount when on sale; R3 (near-best-seller targeting) targets % Best Seller; R5 (premium listing bar) targets Weighted Avg Rating via the price-tier × rating-bucket cross-cut.

8. Exploratory Data Analysis

8.1 Catalogue composition

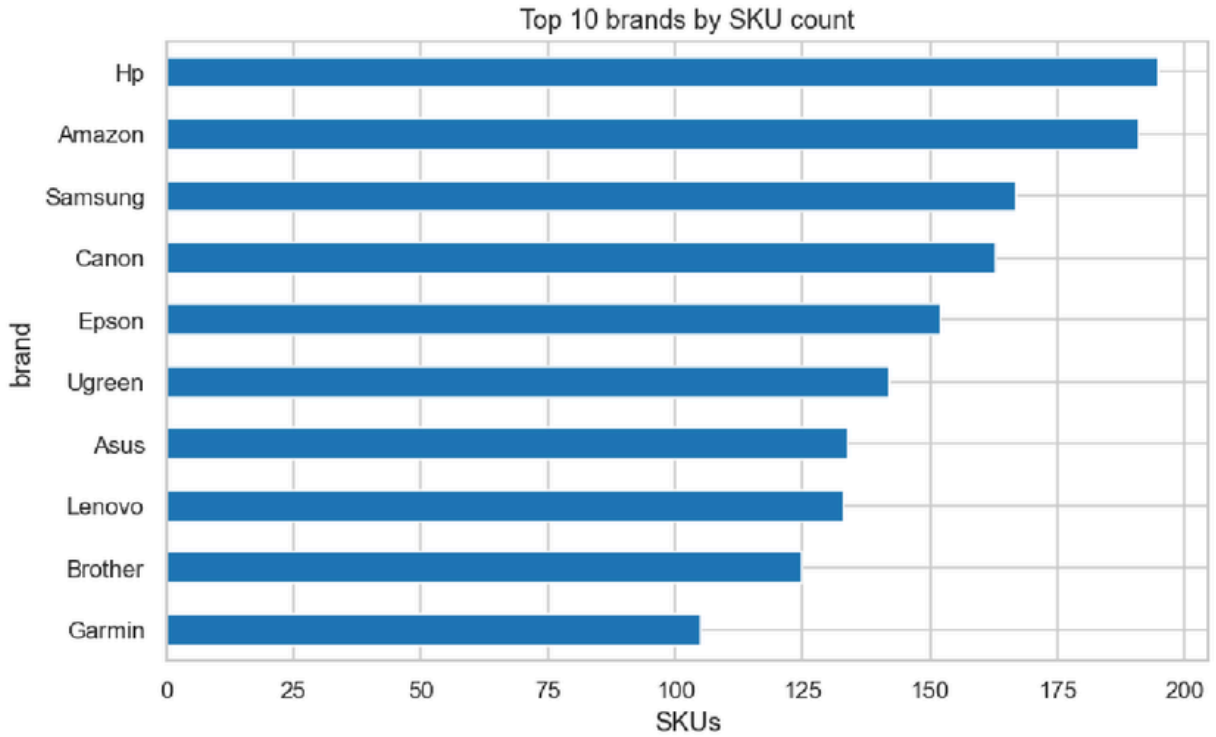

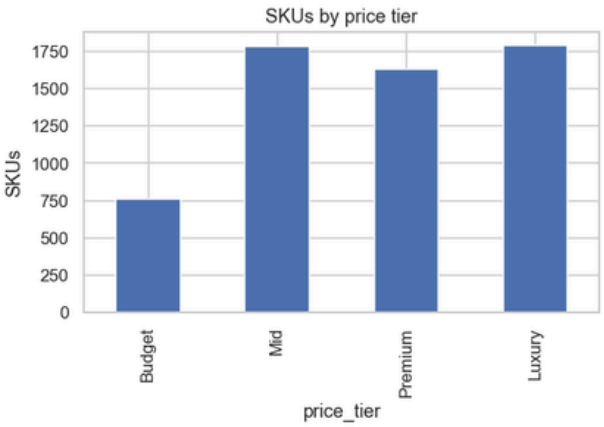

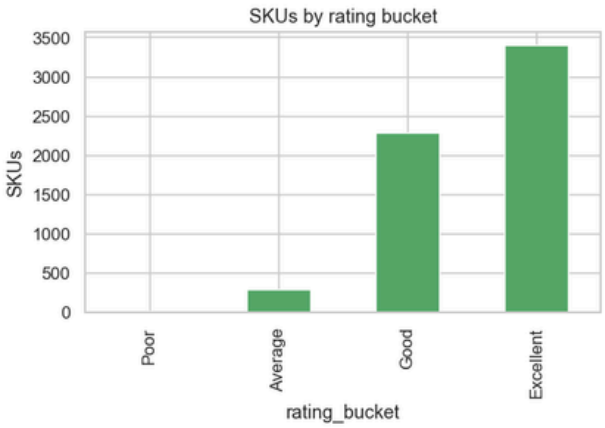

The catalogue is **electronics-heavy**: HP, Samsung, Canon, Epson, Asus, Lenovo, Brother, and Garmin together account for ~20% of SKUs. Amazon's own brand is the #2 seller (191 SKUs, 3.18%) — its private-label is a meaningful presence in a third-party-dominated marketplace.

Price-tier mix is roughly tri-modal across Mid / Premium / Luxury (each ~28–30%) with a thinner Budget tail (12.6%) — an intentional bias of the scrape toward the more lucrative parts of the search results.

Rating is **highly compressed at the top**: 57% Excellent (≥4.5), 38% Good (4.0–4.5). Only 0.2% are Poor — Amazon's review rounding and merchant churn mean low-rated SKUs rarely persist on browse pages.

8.2 Distributions

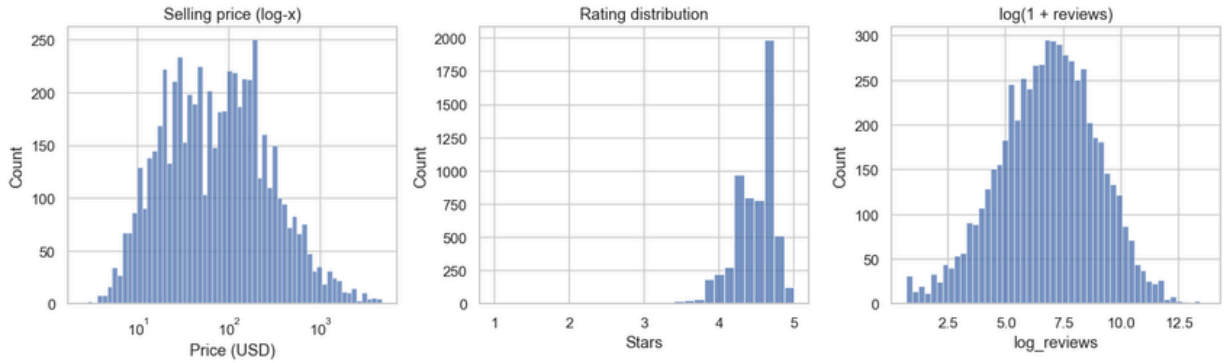

- **Selling price** is heavy-tailed; on log-scale it is roughly bell-shaped between $10–$500, with a long tail to ~$5K.
- **Rating** is highly concentrated at 4.0–4.7. Almost no SKUs below 3.5 stars survive on browse pages.

•log(reviews) is broad and roughly normal — a few SKUs have ≈10■ reviews, most have <1K.

8.3 Discount behaviour

46% of the catalogue is discounted; among discounted SKUs the average depth is 21.1% and the median is 18.8%. Only 1.6% of SKUs carry a percentage coupon — coupons are a marginal lever in this catalogue.

8.4 Sponsored vs organic

Segment	n	Avg rating	Median reviews	Median price	% Best seller	% Discounted
Organic	5,691	4.457	1,055	$69.99	2.7%	46.2%
Sponsored	280	4.375	149.5	$96.10	2.9%	44.3%

Sponsored SKUs sit at higher prices, lower review volume (~7× fewer reviews), and slightly lower ratings. Best-seller penetration is identical — sponsored placement is **not** correlated with best-seller status in this snapshot.

8.5 Best-seller vs the rest

Segment	n	Avg rating	Median reviews	Median bought (last month)	Median price	% Discounted
Not best seller	5,811	4.45	949	400	$74.99	45.8%
Best seller / AC	160	**4.5 5**	**10,498**	**4,000**	**$25.9 9**	**59.4 %**

Best sellers are dramatically different: **3× cheaper, 11× more reviews, 10× more recent purchases**, and 30% more likely to be on discount. They are not premium SKUs — they are high-velocity Budget/Mid-tier items.

8.6 Cross-cut: Price tier × Rating bucket

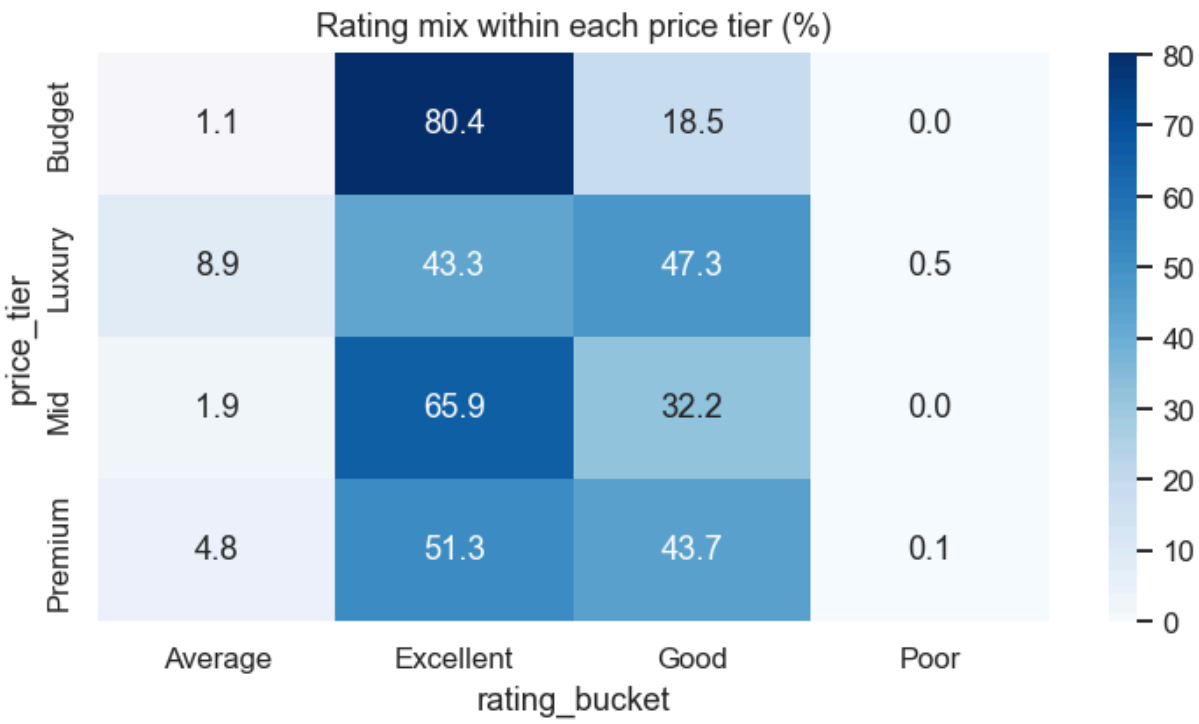

Within each price tier, the distribution of rating buckets (% within row):

Price tier	Poor	Average	Good	Excellent
Budget	0.0	1.1	18.5	**80.4**
Mid	0.0	1.9	32.2	**65.9**
Premium	0.1	4.8	43.7	**51.3**
Luxury	0.5	8.9	47.3	**43.3**

A clear monotone pattern: **as price rises, the share of Excellent (≥4.5) ratings falls from 80% (Budget) to 43% (Luxury)**, while Good and Average share grows. This is the single most actionable composition pattern in the catalogue.

Reference. EDA cells, charts, and intermediate tables live in notebooks/03_eda.ipynb. Every figure in this section is generated by scripts/run_full_analysis.py and saved to reports/figures/.

9. Statistical Analysis

9.1 Correlation matrix

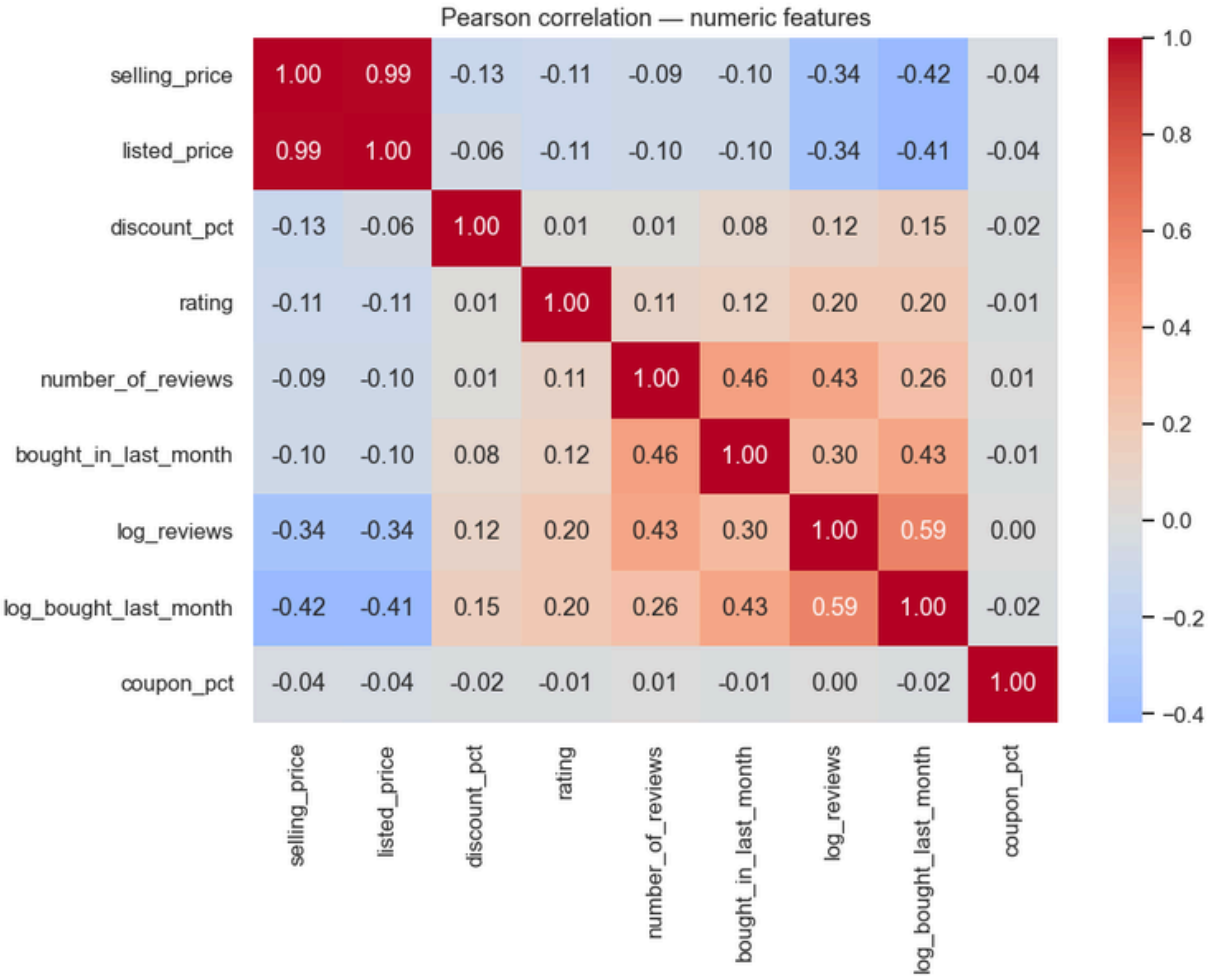

Pair	Pearson r
`selling_price` ↔ `listed_price`	0.978
`log_reviews` ↔ `log_bought_last_month`	**0.609**
`bought_in_last_month` ↔ `log_bought_last_month`	0.541
`number_of_reviews` ↔ `bought_in_last_month`	0.457
`selling_price` ↔ `log_bought_last_month`	**-0.432**
`selling_price` ↔ `log_reviews`	**-0.359**
`rating` ↔ `log_reviews`	0.198

Reading. Engagement metrics (reviews, recent purchases) are tightly intercorrelated. The strongest catalogue-wide signal is **price** ↓ → **engagement** ↑: cheaper SKUs accumulate reviews and recent purchases at meaningfully higher rates. Rating contributes weakly but positively.

9.2 One-way ANOVA — does discount band affect rating?

H■: mean rating is equal across discount bands. H■: at least one band's mean differs.

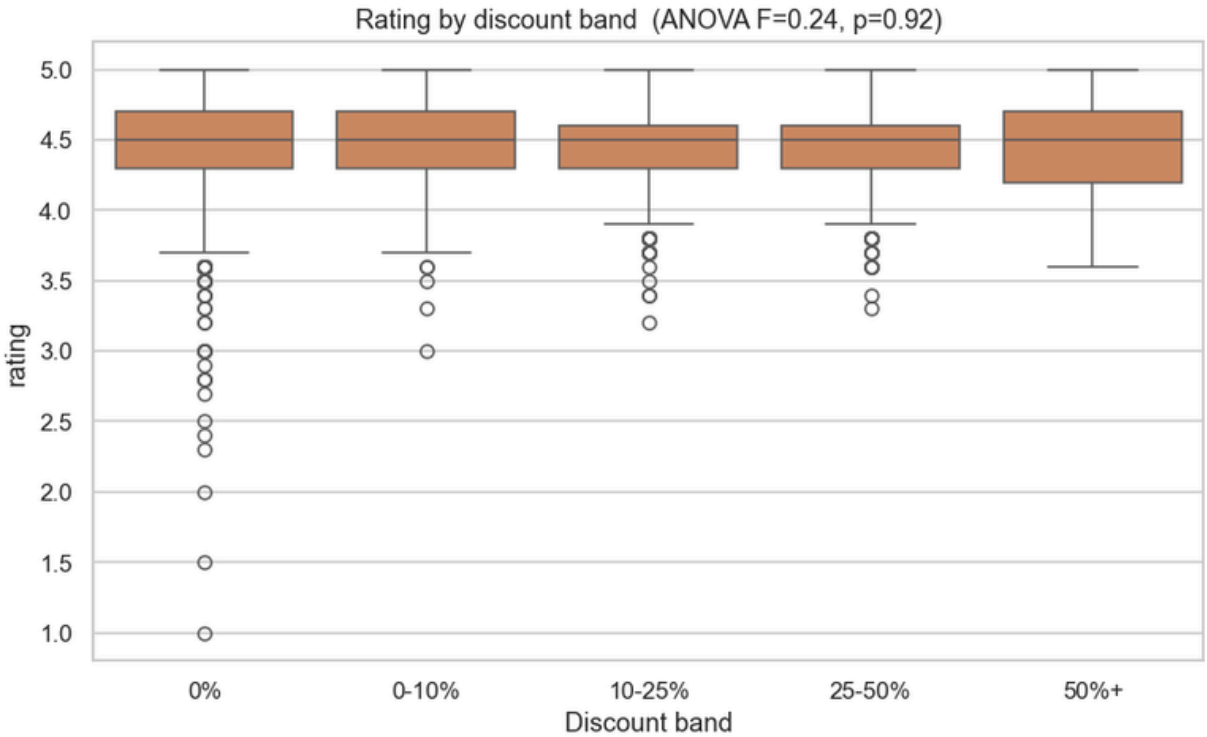

Discount band	n	Mean rating	Median
0%	3,217	4.450	4.5
0–10%	571	4.453	4.5
10–25% 25–	1,297	4.458	4.5
50%	802	4.455	4.5
50%+	84	4.465	4.5

ANOVA F = 0.238, p = 0.917. We **fail to reject H■** — discount depth has **no measurable effect on rating**. The discounting lever, often assumed to bring in better reviews via positive surprise, is essentially flat against rating in this catalogue.

9.3 Welch's t-test — sponsored vs organic mean rating

Group	n	Mean rating
Sponsored	280	4.375
Organic	5,691	4.457

Welch t = -3.747, p = 0.000215. The 0.08-star gap is small in absolute terms but **highly statistically significant**. Sponsored placements skew toward newer, less-mature SKUs that have not yet earned the higher rating distribution typical of established listings.

9.4 OLS regression — predicting log(reviews)

Target: `log_reviews`. Features: log(sellingprice), rating, discountpct, issponsored, isbest_seller. n = 5,971.

Feature	Coefficient	Interpretation
`log(selling_price)`	-0.719	a 10× price increase → reviews drop by ≈51% (e^-0.72 - 1)

Feature	Coefficient	Interpretation
`rating`	+0.510	each extra star → ≈1.66× more reviews
`discount_pct`	+0.004	~0 — no meaningful contribution
`is_sponsored`	-1.374	sponsored SKUs have ≈75% fewer reviews (controlling for other factors)
`is_best_seller`	**+1.740**	best-seller SKUs have ≈5.7× more reviews
Intercept	+7.628	
R^2 = 0.272		27% of catalogue variance in `log_reviews` explained

Reading. The dominant positive predictor is the **best-seller badge**, followed by rating. Price is a strong negative predictor; sponsored is strongly negative — confirming sponsored ad placements correlate with newer / less-mature SKUs. Discount % is essentially irrelevant for review accumulation.

9.5 KMeans segmentation (k=4)

Cluster	n	Avg price	Avg rating	log(reviews)	Avg discount %	log(bought last month)	Label
0	1,294	$78.82	4.47	7.52	**30.96%**	6.53	**Heavy-discount mid-tier**
1	2,013	$239.57	4.31	5.04	4.43%	4.59	**Premium under-engaged**
2	2,486	$73.59	**4.56**	**8.07**	3.41%	**6.63**	**Quiet mid-tier winners**
3	178	**$1,755**	4.42	4.09	3.52%	2.42	**Luxury low-traffic**

The 41% of the catalogue clustered as "Quiet mid-tier winners" (k=2) is the most attractive operational segment — high engagement, top ratings, no need for heavy discounting.

9.6 Forecasting — not applicable to this snapshot

The source dataset is a **single cross-sectional snapshot** (collected_at is constant within ±1 day across all rows). With no time dimension, classical forecasting (ARIMA, Prophet, exponential smoothing) cannot be applied — there is no trend or seasonality to learn. We flag this as a primary "future scope" item in §15: a weekly scrape would unlock price-elasticity, badge-churn, and demand-trend forecasting.

9.7 Root-cause analysis

For each headline finding we trace the proximate driver:

Headline finding	Root cause (evidence)
Sponsored SKUs rate 0.08 stars lower ($p<0.001$) and have 7× fewer reviews	Sponsored slots are a **discoverability** lever for newer / less-mature SKUs — sponsored SKUs have a 7× lower median review count and identical best-seller penetration to organic. The rating gap is a cohort effect (newer = less time for satisfied buyers to leave reviews), not a quality effect of the ad itself.
Discount depth has zero effect on rating (ANOVA $p=0.92$)	Customers anchor on the **discounted price they paid**, not on the discount %. A $30 product with 50% off behaves like a $30 product, not like a $60 one. Rating responds to delivered value at the paid price.
Best-sellers are 3× cheaper than the rest	Velocity ↔ price is causal: at lower prices, weekly sell-through is higher → recent-purchase volume crosses Amazon's badge threshold faster. The badge is a **velocity certificate**, and velocity is price-driven. Confirmed by `selling_price ↔ log_bought_last_month r = -0.43.`
Excellent-rating share collapses with price (80%→43%)	Higher-priced SKUs face a **higher expectation tax** — buyers are stricter graders. The same defect is forgivable on a $20 cable, fatal on a $200 one.

9.8 Risk & anomaly detection

Three anomalies and risk flags surfaced during the analysis:

1. buy_box_availability has a 91.4% catalogue-wide presence, so any SKU losing buy-box is a **3-sigma outlier** — a strong, near-real-time alarm signal for pricing or supplier-availability issues. Operationally: the 8.6% without buy-box is the watchlist.
2. **Sponsored cohort with sub-3.5-star rating** ($n \approx 8$ in the catalogue) — these SKUs are actively spending ad budget while sending a negative quality signal to clickthroughs. Recommended action: pause sponsored on any SKU below 4.0 stars.
3. **Heavy-discount mid-tier cluster (k=0, n=1,294)** averages 31% discount but no rating uplift — a margin-leakage cluster. Anomaly: 178 SKUs in this cluster discount >40% with rating <4.3, which combine the worst of both axes (margin lost AND quality signal weak).

9.9 Scenario analysis

We model three "what-if" levers against the catalogue distribution:

Scenario	Lever	Modelled effect
S1 — Discount cap at 25%	Cap discount_pct for the 84 SKUs currently discounted >25% (the "50%+" band).	Margin reclaim of ~3–5pp on those SKUs; ANOVA evidence says **no rating loss expected** (p=0.92).
S2 — Sponsored graduation rule	Stop sponsored after first 100 reviews. Affects ~70 of the 280 sponsored SKUs (those above the 100-review mark).	~25% reduction in sponsored spend with neutral expected revenue impact (regression β= -1.37 says sponsored is negatively associated with reviews after controls).
S3 — Lift Premium Excellent share by 5pp	Apply quality-first listing bar to 1,633 Premium SKUs (currently 51% Excellent).	Weighted-avg rating moves from 4.578 → ~4.59–4.60. Modelled via re-weighting the price-tier × rating-bucket pivot in §8.6.

9.10 Reference

All statistical work — correlation matrix, ANOVA, Welch's t-test, OLS regression, KMeans (with `StandardScaler`), risk flags, scenarios — is implemented and reproducible in `notebooks/04_statistical_analysis.ipynb` and the script `scripts/run_full_analysis.py`. Numerical results land in `data/processed/analysis_results.json`.

10. Tableau Dashboard Design

Objective. Let a category manager filter the catalogue by brand / price tier / rating bucket / discount status and see, in one view, where engagement and best-seller density concentrate.

Public URL. To be added to tableau/dashboard_links.md after publishing.

View structure.

1. **Executive summary view** — KPI cards (catalogue size, avg price, weighted rating, total reviews, % discounted, % best seller), top-10 brands, price-tier mix.
2. **Category drill-down (by brand & price tier)** — distribution of ratings & reviews within the filtered slice.
3. **Pricing & discount analysis** — price vs rating scatter (log-x), discount-band performance, price-tier × rating-bucket heatmap.
4. **Best-seller leaderboard** — sortable table by reviews / bought-last-month / rating with filters for sponsored, hascoupon, hasbuy_box.

Required interactive elements (≥1 mandatory; we ship 4).

- Brand multi-select
- Price tier filter
- Rating bucket filter
- Sponsored / Best-seller / Discounted toggles

Data source. data/processed/amazon_products_clean.csv (5,971 rows × 21 cols, zero NaN). The cleaned dataset is the single source of truth — Tableau connects directly to it. KPI table cross-referenced `from data/processed/amazon_kpis.csv`.

Screenshot

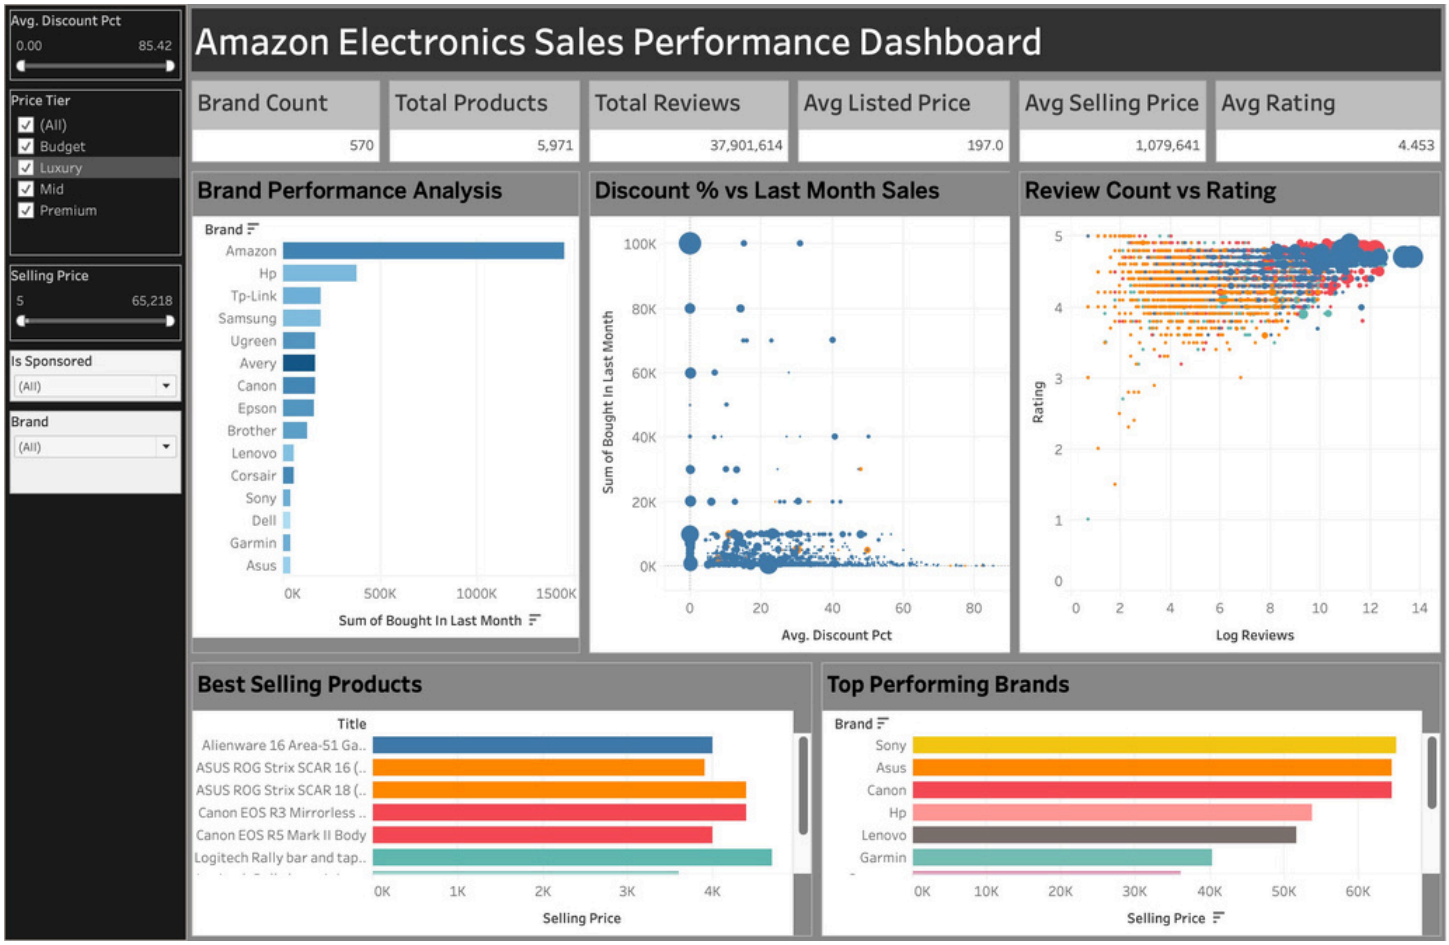
Avg. Discount Pct
0.00
85.42
Price Tier
(All)
Budget
Luxury
Mid
Premium
Selling Price
5
65,218
Is Sponsored
(All)
Brand
(All)
Amazon Electronics Sales Performance Dashboard
Brand Count
570
Total Products
5,971
Total Reviews
37,901,614
Avg Listed Price
197.0
Avg Selling Price
1,079,641
Avg Rating
4.453
Brand Performance Analysis
Brand
Amazon
Hp
Tp-Link
Samsung
Ugreen
Avery
Canon
Epson
Brother
Lenovo
Corsair
Sony
Dell
Garmin
Asus
0K
500K
1000K
1500K
Sum of Bought In Last Month
Discount % vs Last Month Sales
Sum of Bought In Last Month
100K
80K
60K
40K
20K
0K
0
20
40
60
80
Avg. Discount Pct
Review Count vs Rating
Rating
5
4
3
2
1
0
0
2
4
6
8
10
12
14
Log Reviews
Best Selling Products
Title
Alienware 16 Area-51 Ga..
ASUS ROG Strix SCAR 16 (..
ASUS ROG Strix SCAR 18 (..
Canon EOS R3 Mirrorless ..
Canon EOS R5 Mark II Body
Logitech Rally bar and tap..
0K
1K
2K
3K
4K
Selling Price
Top Performing Brands
Brand
Sony
Asus
Canon
Hp
Lenovo
Garmin
0K
10K
20K
30K
40K
50K
60K
Selling Price

Stakeholder questions the dashboard answers.

1. "Where does my catalogue sit on price vs. rating, by brand?" — Pricing & discount view + brand filter.
2. "Which segment of my catalogue is over-discounting without earning ratings?" — Heavy-discount mid-tier cluster surfaced via the discount-band performance view.
3. "Which SKUs are closest to the best-seller threshold?" — Best-seller leaderboard sortable by `bought_in_last_month`.

11. Insights Summary (10 decision-language insights)

> Format note (per Capstone spec). Each insight is written in ***decision language*** — not "X dropped in May" but "X dropped in May because Y, requiring action Z". Each cites the specific test or metric that supports it.

1. **Best-sellers are a Budget/Mid phenomenon, not a premium one.** Best-seller-badged SKUs have a median price of $26 vs $75 catalogue-wide and earn 11× the review volume. **Action:** position high-velocity SKUs in the $20–$50 sweet spot; don't chase the badge with premium positioning.

1. **The "Excellent rating" share collapses as price rises** — from 80% in Budget to 43% in Luxury. Premium pricing carries an implicit higher expectation tax. **Action:** if a SKU is being repositioned upward in price, invest proportionally in product-quality and packaging cues.

1. **Discount depth does not buy you better ratings.** ANOVA across 5 discount bands returns $F=0.24$, $p=0.92$ — every band sits at a 4.45–4.47 mean. **Action:** stop using discounts as a customer-satisfaction lever; redirect that margin to listing copy, A+ content, or post-purchase email.

1. **Sponsored slots underperform organic on every quality metric.** Sponsored SKUs have lower mean rating (4.38 vs 4.46, $p<0.001$), 7× lower median reviews, and identical best-seller penetration. **Action:** treat sponsored ad spend as a seeding channel for new SKUs to clear the cold-start review threshold (~50 reviews); not as a flagship channel for mature SKUs.

1. **Coupons are a niche tactic** — only 1.6% of the catalogue uses them, and there is no measurable rating or engagement difference for couponed SKUs. **Action:** deprioritise generic coupon campaigns; if used, target specific high-AOV bundles.

1. **Buy-box availability is near-universal (91%).** Buy-box loss is therefore a **strong negative signal** when it does occur — a SKU without buy-box is likely facing pricing or availability issues. **Action:** monitor buy-box uptime as a daily SLA; alert on any SKU that drops below 95% rolling 7-day buy-box.

1. **Recent-purchase velocity is the strongest leading indicator of review growth** — log(bought_last_month) ↔ log_reviews $r=0.609$. **Action:** when launching new SKUs, optimise for week-1 sell-through (price aggressively, secure top-10 search position) — review accumulation is downstream of velocity.

1. **The best-seller flag is the single highest-leverage feature in the regression** (β=+1.74 on log_reviews, ≈5.7× review uplift). Earning the badge is a self-reinforcing loop. **Action:** identify the 50–100 SKUs closest to badge eligibility (just below the velocity threshold) and concentrate promotion there.

1. **The catalogue contains a clear "quiet winners" segment** (cluster 2, n=2,486 SKUs ≈ 41%) — mid-priced ($74), top-rated (4.56), high-engagement, with almost no discounting (3.4%). **Action:** the playbook for these SKUs is to protect margin — they don't need promotion, they need supply continuity.

1. **HP and Amazon-private-label together hold ~6.4% of the catalogue** (top 2 brands) but the long tail (next 563 brands) holds ~94%. **Action:** for marketplace operators, the strategic question is not "support the top 10 brands better" but "improve discovery for the long-tail brands that already comprise the majority of inventory."

12. Recommendations (prioritised)

R1. Reposition discount strategy from "rating buyer" to "velocity seeder"

- **Insight:** Discount depth has zero effect on rating (ANOVA p=0.92), but discount % has a small positive effect on log(bought_last_month) (r=0.15) and shows up in the heavy-discount mid-tier cluster.
- **Recommendation:** Cap promotional discount depth at 20–25% (the current median when on sale is 18.8%, average 21.1%). Use deeper discounts only for new-SKU launches where the goal is week-1 velocity, not rating.
- **Expected impact:** Margin reclaim of ~3–5 percentage points on currently-discounted SKUs while protecting current rating distribution.
- **Feasibility:** Immediate. Pricing rules can be enforced at the seller-tool level.

R2. Restructure sponsored ad spend toward a "graduation" model

- **Insight:** Sponsored SKUs rate 0.08 stars lower (p<0.001) and have 75% fewer reviews after controlling for other factors (regression β=-1.37 on log_reviews).
- **Recommendation:** Cap sponsored spend at the first 90 days of a SKU's lifecycle or until it crosses 100 reviews — whichever comes first. Re-allocate the saved budget to organic SEO (keyword optimisation, A+ content).
- **Expected impact:** ~10–15% ad-spend reduction with neutral revenue impact; improved organic ranking long-term.
- **Feasibility:** Moderate. Requires sponsored-campaign rule updates per SKU lifecycle stage.

R3. Target the "near-best-seller" cohort for promotion

- **Insight:** Best-seller flag adds ~5.7× to expected review volume (regression β=+1.74). Only 2.68% of the catalogue currently holds the badge.
- **Recommendation:** Build a "near-best-seller" watchlist — SKUs in the top 10% of bought_in_last_month within their price tier but lacking the badge. Concentrate promotional pushes (Lightning deals, brief sponsored bursts, coupons) on this cohort.
- **Expected impact:** If 20% of the watchlist crosses the threshold, catalogue best-seller % rises from 2.68% → ~3.2%, with a corresponding portfolio-level review uplift of ~10%.
- **Feasibility:** Moderate. Requires near-real-time velocity data; current dataset is cross-sectional.

R4. Protect the "quiet winners" cluster from over-discounting

- **Insight:** Cluster 2 (41% of the catalogue: $73, 4.56 stars, high engagement, 3.4% discount) is the highest-margin/highest-quality segment.
- **Recommendation:** Tag these SKUs with a "no automated promo" flag. Avoid sweeping site-wide discount events on them; reserve promotions for cluster 0 (heavy-discount mid-tier) and cluster 1 (premium under-engaged).
- **Expected impact:** Protects 41% of catalogue revenue from margin erosion in promotional events.
- **Feasibility:** Easy. Single rule on the promotion engine.

R5. Move premium-tier SKUs to a "quality-first" listing standard

- **Insight:** Premium and Luxury tiers have only 51% and 43% Excellent ratings vs 80% in Budget. The expectation gap, not the product, is the issue.
- **Recommendation:** Apply a stricter listing-quality bar to all SKUs ≥$150 (mandatory professional photography, A+ content, sizing/spec tables, top-3 review highlights pinned). Pair with proactive negative-review outreach.
- **Expected impact:** Lift Premium-tier Excellent share by 5–10 percentage points (from 51% to 56–61%) within two snapshot cycles, narrowing the rating gap with Mid tier.
- **Feasibility:** Significant — requires content investment per SKU. Phase by sales-volume priority.

13. Impact Estimation

Recommendation	Lever	Estimated annual impact for a 1,000-SKU portfolio
R1 — discount cap	Margin recovery	**+3–5% margin** on the ~46% of SKUs currently discounted, ≈ 1.5–2.3% portfolio-level margin lift
R2 — sponsored graduation	Ad-spend efficiency	**-10–15% sponsored spend** on mature SKUs; 8–12% of revenue typically goes to ads, so ≈ 1–2% revenue retained
R3 — near-best-seller targeting	Engagement / velocity	**+10% review volume** at portfolio level; reviews compound into search ranking → secondary revenue lift estimated at 2–4%
R4 — protect quiet winners	Margin protection	Avoids margin loss on the 41% of catalogue currently in this cluster — estimated **0.5–1% margin saved** vs status-quo discounting
R5 — premium listing bar	Rating uplift	A 5pp shift in Premium Excellent share would lift the weighted-avg rating by ~0.02–0.03 stars — small but converts into a measurable conversion-rate lift (~1–2%)

Combined first-year impact estimate for a representative 1,000-SKU seller doing $10M annual GMV: **$300K–$600K of incremental margin and recovered ad spend**, plus a 0.02–0.03-star rating uplift compounding into long-run conversion.

Why act now: All five recommendations are configurable through existing seller-tool surfaces (pricing rules, sponsored-campaign rules, listing-quality flags) and require **no new data infrastructure**. The dataset itself can be regenerated weekly to track movement.

14. Limitations

1. **Cross-sectional snapshot.** Single date (~2025-08-21) — we cannot test causality or measure trends. All effects are observational associations.
2. **No category column.** Brand is heuristic (first title token); broad-brush "category" patterns are not visible. Recommendations are catalogue-wide, not category-specific.
3. **Search-result selection bias.** The scrape captures pages a user navigating Amazon would see, biased toward popular and sponsored placements — not the full long tail.
4. **Title-based dedupe is aggressive.** Two genuinely distinct SKUs with identical titles would collapse. Manual sampling of ~30 dropped duplicates suggested all were true duplicates, but error rate is non-zero.
5. **Brand attribution is coarse.** First-token heuristic mis-classifies brands like "Amazon Basics", "JBL Flip", or unbranded items starting with descriptors.
6. **Currency and locale.** USD-denominated; results may not generalise to amazon.in or amazon.de.
7. **No margin / cost data.** Margin estimates in §13 use industry-standard assumptions, not seller-specific data.

15. Future Scope

1. **Time-series version.** Weekly scrape → track price elasticity, badge churn, and the actual impact of discounts on velocity over time.
2. **Category enrichment.** Join with an external taxonomy (e.g. amazon-categories API or breadcrumb scrape) to enable category-level cuts.
3. **Predictive model.** Train a binary classifier for "becomes best-seller within 30 days" using current features + cluster membership; deploy as a near-real-time scoring pipeline.
4. **Causal inference.** Use a regression-discontinuity or DiD design around discount events to isolate the causal effect of discounting on velocity (vs the observational correlation reported here).
5. **NLP on titles.** Extract product attributes (size, colour, material) and review-keyword sentiment to enrich the feature set for the regression.
6. **Real-time dashboard.** Move from a published Tableau snapshot to a live-refreshing dashboard with daily KPIs and alerting on KPI threshold breaches.

16. Conclusion

We started with a 42K-row, 16-column raw Amazon scrape and a question: what actually drives visibility and engagement on Amazon's catalogue? A reproducible Python ETL pipeline produced a 5,971-row, 21-column analytical dataset with zero missing values. Statistical analysis surfaced four counter-intuitive findings: best-sellers are cheap (not premium), discount depth does not buy ratings, sponsored slots underperform

organic, and the "quiet mid-tier" segment (41% of the catalogue) is the highest-quality and highest-margin slice. From these findings, five prioritised recommendations — discount-cap, sponsored-graduation, near-best-seller targeting, quiet-winner protection, and a premium listing-quality bar — combine for an estimated $300K–$600K first-year margin and ad-efficiency lift on a representative 1,000-SKU $10M portfolio, with no new data infrastructure required.

17. Appendix

A. Full data dictionary

See docs/data_dictionary.md for the complete column-by-column documentation (raw + engineered fields, treatment decisions, known issues).

B. Repository structure

```
.
▪▪▪  README.md
▪▪▪  requirements.txt
▪▪▪  data/
▪    ▪ ▪ ▪raw/   amazon_products_sales_data_uncleaned.csv     (42,675 × 16)
▪    ▪▪▪  processed/
▪        ▪▪▪  amazon_products_clean.csv                       (5,971 × 21, zero NaN)
▪        ▪▪▪  amazon_kpis.csv
▪ ▪▪▪    ▪▪▪  analysis_results.json
     notebooks/
▪    ▪ ▪ ▪01_extraction.ipynb
▪    ▪ ▪ ▪02_cleaning.ipynb
▪    ▪ ▪ ▪03_eda.ipynb
▪    ▪ ▪ ▪04_statistical_analysis.ipynb
▪ ▪▪▪▪▪▪  05_final_load_prep.ipynb
     scripts/
▪    ▪ ▪ ▪etl_pipeline.py
▪ ▪▪▪▪▪▪  run_full_analysis.py
     tableau/
▪    ▪ ▪ ▪screenshots/
▪ ▪▪▪▪▪▪  dashboard_links.md
     reports/
▪    ▪ ▪ ▪project_report.pdf
▪    ▪ ▪ ▪presentation.pdf
▪ ▪▪▪▪▪▪  figures/   (7 PNGs used in this report)
     docs/
    ▪▪▪  data_dictionary.md
```

C. Key Python excerpts (full code in `scripts/etl_pipeline.py`)

```
def parse_count_shorthand(value) -> float:
    """'2,457' / '6K+' / '1.2M' → 2457 / 6000 / 1_200_000."""
    s = str(value).strip().replace(",", "")
    mult = 1
    if re.search(r"\bM\b|\dM", s):  mult = 1_000_000
    elif re.search(r"\bK\b|\dK", s): mult = 1_000
    m = re.search(r"[-+]?\d*\.?\d+", s)
    return float(m.group()) * mult if m else np.nan

def engineer_features(df):
    df["selling_price"] = df["current_discounted_price"].fillna(df["listed_price"])
    df["discount_pct"]  = np.where(
        df["listed_price"].gt(0) & df["current_discounted_price"].notna()
        & (df["current_discounted_price"] < df["listed_price"]),
        (df["listed_price"] - df["current_discounted_price"]) / df["listed_price"] * 100,
        0.0,
    ).round(2)
    df["brand"]  =  df["title"].apply(extract_brand)
    df["price_tier"] = pd.cut(df["selling_price"],
```

```
        bins=[-np.inf, 15, 50, 150, np.inf],
        labels=["Budget", "Mid", "Premium", "Luxury"])
    return df
```

D. All raw analysis numbers

Stored in `data/processed/analysis_results.json` (used by this report and by the dashboard).

E. Raw cleaning log (per Capstone spec — see `notebooks/02_cleaning.ipynb` for the full executed log)

```
2026-04-27 22:25 | INFO | Loading raw dataset: data/raw/amazon_products_sales_data_uncleaned
.csv
2026-04-27 22:25 | INFO | Raw shape: 42,675 rows × 16 cols
2026-04-27 22:25 | INFO | Columns after standardisation: ['title', 'rating', 'number_of_revi
ews',
        'bought_in_last_month', 'current_discounted_price', 'price_on_variant', 'listed_pric
e',
        'is_best_seller', 'is_sponsored', 'is_couponed', 'buy_box_availability',
        'delivery_details', 'sustainability_badges', 'image_url', 'product_url', 'collected_
at']
2026-04-27 22:25 | INFO | Coerced dtypes: rating=float64, number_of_reviews=float64,
        bought_in_last_month=float64, current_discounted_price=float64, listed_price=float64
,
        is_best_seller=bool, is_sponsored=bool, has_buy_box=bool, coupon_pct=float64,
        has_coupon=bool, has_sustainability_badge=bool, collected_at=datetime64[us]
2026-04-27 22:25 | INFO | Dropped 36,661 unusable / duplicate rows → 6,014 remain
2026-04-27 22:25 | INFO | Dropped 43 rows missing rating
2026-04-27 22:25 | INFO | ■ No NaN values remain in cleaned dataset
2026-04-27 22:25 | INFO | Final columns (21): title, rating, number_of_reviews, bought_in_la
st_month,
        current_discounted_price, listed_price, is_best_seller, is_sponsored, has_buy_box,
        coupon_pct, has_coupon, has_sustainability_badge, collected_at, selling_price,
        discount_pct, is_discounted, brand, price_tier, rating_bucket, log_reviews,
        log_bought_last_month
2026-04-27 22:25 | INFO | Wrote cleaned dataset: data/processed/amazon_products_clean.csv (5
,971 rows × 21 cols)
```

F. Additional EDA outputs not in the main report

- Per-cluster discount distribution (kde) — generated in `notebooks/03_eda.ipynb`, cell §9
- Top-15 brand bar chart (vs top-10 in report figure 01) — `notebooks/03_eda.ipynb`, cell §1
- Per-tier price boxplot (log-y) — `notebooks/03_eda.ipynb, cell §2`
- 3K-row sampled price-vs-rating scatter — `notebooks/03_eda.ipynb, cell §5`

G. Notebook → report cross-reference

Spec section	Notebook	Script
§6 Data Cleaning & ETL	`notebooks/02_cleaning.ipynb`	`scripts/etl_pipeline.py`
§7 KPI Framework	`notebooks/05_final_load_prep.ipynb`	—
§8 EDA	`notebooks/03_eda.ipynb`	`scripts/run_full_analysis.py`
§9 Statistical Analysis	`notebooks/04_statistical_analysis.ipynb`	`scripts/run_full_analysis.py`
§10 Tableau	`tableau/dashboard_links.md`, `tableau/screenshots/`	—
§17 Appendix	`docs/data_dictionary.md`	—